INTERNAL PRODUCT DOCUMENT

Guest CRM

Project Structure, Use Cases, Roadmap and Product Logic

Version: 2.0 — June 2026

For: Product & Engineering Team

Tool: Chekin · Guest CRM

1 Executive summary

A guest intelligence and revenue-activation layer

This document defines the product logic and roadmap for the Chekin **Guest CRM**: a guest intelligence and revenue-activation module that lets Chekin's B2B clients understand their guests, build audiences in plain language, create campaigns with an AI agent, attach App Sales deals, trigger the existing Automation Engine and measure the revenue generated.

The opportunity is not to build a generic CRM. It is to build an **agentic CRM** that already knows the reservation context, the guest profile, the property, the timing of the trip and the available commercial inventory, and turns that intelligence into personalised communication and revenue per reservation — without breaking guest trust.

> **North Star** Turn every guest into an understood profile, every profile into a relevant audience, and every audience into a revenue opportunity that still feels helpful, trusted and guest-first.

What this product should achieve

- ✓ Give the client a single, intelligent view of each guest (Guest 360).
- ✓ Let non-technical staff build audiences by typing, not by configuring rules.
- ✓ Let an agent draft campaigns end-to-end as structured, editable objects.
- ✓ Monetize through relevant App Sales deals attached at the right moment.
- ✓ Reuse the existing Automation Engine and Unified Inbox, not replace them.
- ✓ Keep consent, suppression and human approval as first-class guardrails.
- ✓ **Score every guest automatically** with a transparent, rule-based Opportunity Score (0–100) so operators know at a glance who to prioritise.
- ✓ **Tag guests with CRM decisions** (Whale, VIP, Do Not Contact…) so the whole team shares a common vocabulary — no ML, no black boxes.

2 Product vision and positioning

A guest revenue activation layer, not a database

The product should be positioned as a **guest revenue activation layer**, not as a database. It should feel like a smart operator who already knows the stay context and can anticipate what to offer, to whom, on which channel and when — while the client stays in control of every send.

Proposed product names

Name	Why it works
Guest CRM	Clear, broad, instantly understood by clients.
Guest Intelligence	Emphasises the 360 profile and insight angle.
Revenue CRM	Strong for the monetization narrative.
Guest Journey CRM	Ties pre-, in- and post-stay together.
Stay Marketing	Simple and action-oriented.

> **Recommended positioning** "Know your guests, sell more, stay trusted." The message is simple for the client and commercially powerful for Chekin: Chekin is no longer only managing the stay; Chekin is helping monetize the relationship around it.

3 Why Chekin has an unfair data advantage

Structured stay context from day one

Most CRMs start with a cold contact list. Chekin starts with structured, high-intent stay context: the guest has booked, has a property, has dates, has a party composition, is checking in, and is already communicating through the Unified Inbox. This is a far stronger targeting and recommendation surface than a generic marketing CRM.

Data layer	Examples
Reservation context	Property and location; check-in / check-out dates; arrival and departure timing; nights, channel, booking value and spend tier.
Guest and party context	Language and country; family / couple / business inference; children or companions; lifecycle stage; repeat and VIP signals.
Commercial inventory	Own and third-party App Sales deals; price, margin and commission; per-property availability and eligibility.
Behavioural and conversational	Inbox intents (e.g. asked about parking); clicked and dismissed offers; check-in completion; past purchases and stays; consent state.
Segmentation signals	Operator-assigned CRM tags (Whale, VIP, Preferred…); computed Opportunity Score (0–100); negative flags (Do Not Contact, Unwanted Guest).

Differentiation vs generic CRMs

- ✓ Generic CRMs ask you to import and segment cold data; Chekin already holds live stay context.
- ✓ Email tools blast lists; Chekin can trigger on real stay events via the Automation Engine.
- ✓ Marketplaces show offers; Chekin can place the right deal at the right moment in the journey.
- ✓ Other CRMs guess consent; Chekin captures it at check-in and enforces it at send time.
- ✓ Other CRMs rely on manual tagging; Chekin auto-scores every guest with a transparent 9-signal algorithm.

4 Product architecture

A modular system connected through one guest model

The Guest CRM is designed as a modular system. Guest intelligence, segmentation, campaign creation, deal recommendation and automation are separable enough to evolve independently, but connected through a single guest model and a shared agent orchestration layer.

Module	Responsibility
Guest context builder	Unifies reservation, identity, preference, consent and conversation data into the Guest 360 profile.
Opportunity scoring engine	Computes a transparent 0–100 score for every guest from 9 rule-based signals. Score is always explainable; each guest sees up to 5 plain-language reasons.
Guest tag manager	Lets operators assign and remove 12 CRM decision tags (Whale, VIP, Preferred, Repeat Guest, Family, Business, Couple, High App Sales Potential, Needs Review, Do Not Contact, Do Not Promote, Unwanted Guest). Negative tags require a reason note.
Segment builder	Turns a natural-language request into structured, editable audience rules with a live size preview. Understands score thresholds and tag filters.
Campaign studio (agent)	Drafts objective, audience, deals, copy, variants, automation and revenue estimate as one editable object.
Deal recommendation engine	Ranks own and third-party deals by fit with the audience and historical conversion.
Consent & guardrail layer	Runs a consent check on every campaign, excludes unconsented guests and blocks unsafe activation.
Automation bridge	Binds approved campaigns to existing Automation Engine triggers; supports test runs.
Analytics layer	Tracks sends, opens, clicks, conversions and revenue; summarises performance with the agent.

> **Design principle** Each AI output is a **structured, editable object**, never free text alone. The agent drafts; the user reviews and approves; the engine executes. This keeps the experience agentic but controlled and explainable.

5 UX flow and navigation

Eight screens, one coherent experience

The Guest CRM lives inside the existing Chekin dashboard as a product area (sidebar item between Upselling and Documents), reusing the design system, sidebar, Vela agent panel and layout conventions. No separate visual identity is introduced.

- ✓ Guest CRM dashboard with an agentic command box, AI suggestions, and **segmentation stat widgets** (high-opportunity guests, Whales, VIPs, Repeat Guests, App Sales targets and excluded guests).
- ✓ Guests list with filters and a Guest 360 profile drawer. Columns include a live **Opportunity Score pill** and **CRM tag chips**. Quick-filter chips: 🐋 Whale, ⭐ VIP, 🔁 Repeat, 📈 High App Sales, ⚠ Excluded.
- ✓ Natural-language segment builder with a live audience preview. Accepts score thresholds and tag filters alongside reservation rules.
- ✓ Agentic campaign studio (six guided steps).
- ✓ Campaigns list and campaign detail with approval and analytics.
- ✓ Deals catalog (own and third-party).
- ✓ Automation bindings and simulated test runs.
- ✓ Consent and preference settings with a suppression list and audit trail.

> **Current HTML prototype alignment** The current prototype already covers all of the above: dashboard (with segmentation widgets and AI insight strip), guests + profile drawer (with score panel and inline tag manager), segment builder (with score and tag rules), six-step campaign studio, campaign detail, deals, automation and consent — all running on a typed mock data / service / agent layer that is ready to be swapped for real APIs.

6 Agentic campaign studio concept

One prompt, one complete campaign

The campaign studio is the emotional and commercial hook. A user types a goal in plain language and the agent returns a complete, structured campaign that can be inspected and edited at every step. It should feel like delegating to a competent marketer who shows their work.

From a single prompt the agent generates the objective, the audience rules and estimated size, the recommended deals with a fit explanation, the channel, the timing trigger, the message copy and an A/B variant, the consent check, and the revenue estimate — each as an editable field with a visible "Why?".

Example campaign prompts

🏝️

Pre-arrival upsell

"Sell late checkout to Sunday departures" → audience: checkoutDay = Sunday; deal: late checkout €30; channel: email 24h before.

👪

Family cross-sell

"Reactivate past summer family guests" → audience: tags = family; lastStayAt last summer; nextStayAt empty.

🐳

Whale outreach

"Target whale guests with no future reservation" → audience: segmentTag = whale; nextStayAt is empty. Score ≥ 75 auto-added.

⭐

VIP retention

"Email VIPs with Opportunity Score above 70" → audience: segmentTag = vip; opportunityScore ≥ 70; consent = email.

Consent always wins Regardless of tags or score, the consent guardrail runs last. Guests flagged Do Not Contact, Unwanted Guest or Do Not Promote are always excluded before any send — even if they hold a Whale or VIP tag.

7 Segment model and natural-language audiences

Type a sentence, get a structured rule set

Audiences are the input to every campaign. Instead of a rule-builder maze, the user describes the audience in a sentence and the agent parses it into structured rules with a live eligible-size preview after consent.

Example prompt	Generated rules (structured)
Families arriving this weekend with children	checkInDate between today..+3d; children > 0
Guests who have not completed check-in	checkinStatus = abandoned; checkInDate in next 7d
Past summer families with no future booking	stayed last summer; tags contains family; nextStay is empty
Guests who asked about parking, arriving in 72h	intent = parking_inquiry; checkInDate in 72h
Whale guests with no future reservation	segmentTag = whale; nextStayAt is empty
Guests with Opportunity Score above 70	opportunityScore >= 70; excludeNegative = true
High App Sales potential guests arriving soon	segmentTag = high_app_sales_potential; checkInDate in next 14d
Exclude Unwanted Guests from all sends	segmentTag not in [do_not_contact, unwanted_guest, do_not_promote]
Repeat guests with email consent	segmentTag = repeat_guest; consent = email
VIP guests who stayed last year	segmentTag = vip; lastStayAt in 2025

Every rule is editable as a chip, the audience size recomputes live, and consent exclusions are shown explicitly. Segments can be saved and reused, or sent straight into the campaign studio.

8 Guest Segmentation MVP

Opportunity Score, Guest Tags and the classification layer

The Guest Segmentation MVP adds a transparent classification layer on top of every guest profile. It gives operators two tools: a **rule-based Opportunity Score** (0–100) that is always explainable, and a set of **12 CRM decision tags** that encode the team's shared vocabulary about each guest.

> **Core principle** No machine learning. No predictive models. No hidden rules. Everything is understandable by a property manager in seconds. Each score shows its top reasons in plain language; each tag is assigned by a human and can be removed at any time.

Opportunity Score (0–100)

The score is computed in real time from nine additive signals and one negative penalty. It represents how commercially attractive a guest is right now — not a prediction, but a snapshot of current context.

Signal	Points	Condition
Upcoming reservation	+25	Check-in date within the next 90 days
Repeat guest	+15	At least one previously completed stay
App Sales history	+15	At least one App Sales purchase on record
High booking value	+10	Total booking value ≥ €2,000
Email marketing consent	+10	Email marketing consent flag is active
Complete profile	+10	Profile completeness ≥ 75%
WhatsApp consent	+5	WhatsApp consent flag is active
Completed check-in	+5	Online check-in has been completed
Currently in-house	+5	Guest is currently at the property

Signal	Points	Condition
Global unsubscribe	−20	Guest has globally unsubscribed from all comms

Guests bearing a negative tag (Do Not Contact, Unwanted Guest or Do Not Promote) have their score forced to 0 and are excluded from all campaign audiences, regardless of their underlying signal values.

🟢

High (75–100)

Prime targets. Multiple positive signals active. Prioritise for personalised outreach.

🟡

Medium (40–74)

Good candidates. Some signals missing. Nurture with relevant offers.

🔴

Low (0–39)

Thin context. Limited history or consent. Avoid unsolicited contact.

Score in the UI

- ✓ **Guests table:** compact score pill with a colour-coded mini-bar in the Score column.
- ✓ **Guest drawer:** full score panel at the top showing the numeric value, tier label and up to 5 plain-language reasons (e.g. "Upcoming reservation this week · Repeat guest · Email consent active").
- ✓ **Dashboard:** segmentation stat widgets — total high-opportunity guests (score ≥ 75), Whale count, VIP count, Repeat Guest count, High App Sales Potential count, and excluded guests count.
- ✓ **Segment builder:** accepts `opportunityScore >= N` as a first-class filter rule.

Guest Tags — the 12-tag vocabulary

Tags encode CRM decisions made by operators. Unlike descriptive tags derived from booking data (family, couple, business), CRM tags are intentional — they are assigned and removed manually, require a note when negative, and travel with the guest profile across stays.

Tag	Category	When to use
🐳 Whale	Commercial	Highest-value guests; large cumulative revenue or spend potential.

Tag	Category	When to use
⭐ VIP	Commercial	Guests who receive elevated service or priority treatment.
🏅 Preferred	Commercial	Guests with preferred status — loyalty or operator discretion.
🔁 Repeat Guest	Loyalty	Guests who have completed more than one stay at the property.
👪 Family	Profile	Guest travelling with children or a family group.
💼 Business	Profile	Guest travelling for work; invoice or co-working likely relevant.
❤️ Couple	Profile	Guest travelling as a couple; wellness and dining upsells relevant.
📈 High App Sales Potential	Commercial	Guest likely to purchase App Sales deals based on profile or history.
⚠️ Needs Review	Operational	Profile or behaviour needs operator attention before next send.
🚫 Do Not Contact	Negative	Guest requested complete suppression. Requires a written reason.
🔇 Do Not Promote	Negative	Guest opted out of promotional sends only. Requires a written reason.
❌ Unwanted Guest	Negative	Guest is not welcome back. Requires a written reason.

Negative tags require a reason Do Not Contact, Do Not Promote and Unwanted Guest cannot be assigned without a written note. This creates an audit trail and ensures the tag is never applied carelessly. Guests with any of these three tags are automatically excluded from all campaign audiences.

Tag manager in the guest drawer

Operators can add and remove tags directly from the guest profile drawer without leaving the guest view. A tag picker popover shows all 12 tags with colour coding. Assigned tags appear as chips below the score panel with a remove button. All changes take effect immediately on the next segment or campaign preview.

Dashboard segmentation widgets

High Opportunity Guests with score ≥ 75 — prime targets for any campaign.	**Whales** Guests tagged 🐳 Whale — highest-value segment.	**VIPs** Guests tagged ⭐ VIP — elevated service expectation.
Repeat Guests Guests tagged 🔁 Repeat — proven loyalty signal.	**High App Sales** Guests tagged 📈 High App Sales Potential.	**Excluded** Guests with a negative tag — suppressed from all sends.

Below the stat widgets, an AI insight strip surfaces one contextual recommendation, e.g. "*3 Whale guests have an upcoming stay this week — consider a personalised pre-arrival offer.*" The strip links directly to Campaign Studio.

AI Agent segmentation support

The CRM agent understands segmentation prompts natively. It can parse score thresholds, tag names and exclusion requests from natural language and translate them into structured segment rules.

💬

Score-based prompt

"Show me guests with Opportunity Score above 70" → opportunityScore ≥ 70; excludeNegative = true

🐳

Tag-based prompt

"Whale guests with no future reservation" → segmentTag = whale; nextStayAt is empty

🚫

Exclusion prompt

📈

Combined prompt

"Exclude Unwanted Guests from all sends" → segmentTag not in [do_not_contact, unwanted_guest, do_not_promote]

"High App Sales potential guests arriving this month" → segmentTag = high_app_sales_potential; checkInDate in next 30d

9 Use case roadmap

The same engine, a growing set of revenue moments

The same engine powers a growing set of revenue and retention use cases across the journey.

Use case	Trigger and offer
Pre-arrival upsell	Days before arrival → early check-in, parking, transfer, crib.
Check-in recovery	Check-in abandoned → operational nudge to finish before arrival.
In-stay cross-sell	During stay → breakfast, spa, late checkout, experiences.
Departure upsell	Before checkout → late checkout for relaxed departures.
Win-back / rebooking	Past guests, no future booking → seasonal reactivation.
Intent-based offers	Inbox intent (parking, cars) → targeted, fast WhatsApp offer.
Whale / VIP personalisation	Operator-tagged high-value guests → bespoke pre-arrival or in-stay outreach.
High App Sales outreach	Guests tagged High App Sales Potential + upcoming stay → curated deal bundle.

10 Deals and monetization taxonomy

Own and third-party deals through the App Sales layer

Monetization runs through the existing App Sales layer. Deals are own (margin) or third-party (commission), scoped per property, with conversion and average-revenue history that feeds recommendation.

Deal category	Examples and economics
Convenience (own)	Early check-in, late checkout, baby kit — high margin, high conversion.
Transport	Parking (own), airport transfer and bike rental (third-party commission).
Dining	Breakfast delivery (own); family or partner restaurant sets (commission).
Wellness & experiences	Spa passes and experiences (third-party commission, high average value).
Business	Co-working day passes for business travellers.

11 Recommendation and monetization logic

Score, fit, explain — never just rank

For a given audience, the agent scores each active deal by tag overlap with the audience signals and by historical conversion rate, then explains the fit. Own deals are weighted for margin, third-party for commission and average revenue.

- ✓ Score = base conversion + audience-fit boosts (family, couple, business, intent, timing).
- ✓ Each recommendation carries a one-line "Why?" so the choice is explainable.
- ✓ Revenue estimate = eligible audience × expected conversion × deal price.
- ✓ Offers are subtle and contextual — the right deal at the right journey moment, not advertising.
- ✓ Guests with a low Opportunity Score or a negative tag are excluded from deal recommendations automatically.

12 Engineering implementation model

Typed mock layer ready to swap for real APIs

The current HTML prototype is built with vanilla JavaScript (ES5), plain HTML5 and plain CSS3 — zero external runtime dependencies. This keeps the prototype fast to iterate and easy to inspect.

Layer	Role
data/types.js	Typed constants: lifecycle stages, consent types, segment tag vocabulary.
data/fixtures.js	29 deterministic guest records with segmentation tags, scores and consent states pre-applied.
services/crm-service.js	Pure-function service layer: scoring engine, tag management, segment filters, segmentation metrics. All logic is synchronous and test-friendly.
agent/agent.js	Natural-language parser → structured segment rules. Understands score thresholds and all 12 tag names.
ds/_crm.js	UI component library: score pills, score panels, tag chips, tag picker, drawer, toast, campaign preview.
ds/_crm.css	Design tokens and component styles. Uses CSS variables from the Chekin design system.

> **API-ready pattern** All data access goes through `window.CRMService`. Swapping the fixture layer for real API calls requires changes in one file only. The scoring engine, agent parser and UI components are fully decoupled from the data source.

13 Analytics and success metrics

What good looks like

The product measures success across three levels: operator adoption, campaign performance and revenue impact.

Adoption

% properties with at least one active campaign; % guests with at least one CRM tag assigned.

Campaign

Open rate, click rate, conversion rate and revenue per campaign. Tracked per channel.

Revenue

Total revenue attributed to Guest CRM campaigns. Revenue per eligible guest per month.

Segmentation-specific metrics

- ✓ % guests with a computed Opportunity Score ≥ 75 who received a targeted campaign.
- ✓ Conversion rate difference: campaigns targeting score ≥ 75 vs. unscored bulk campaigns.
- ✓ % Whale and VIP guests with at least one personalised contact per stay.
- ✓ % campaigns that excluded at least one negative-tagged guest (guardrail effectiveness).

14 Consent, trust and guardrails

Consent is not a checkbox — it is an architecture decision

Every campaign runs a consent check before activation. Guests without the required consent for the selected channel are excluded from the eligible audience, and the exclusion count is shown explicitly in the campaign studio.

Guardrail	Enforcement point
Email consent required for email sends	Segment builder + campaign studio pre-check
WhatsApp consent required for WhatsApp sends	Segment builder + campaign studio pre-check
Global unsubscribe respected	Scoring engine (–20 points) + hard exclusion at send
Do Not Contact / Unwanted Guest / Do Not Promote	Tag manager + hard exclusion from all campaign audiences
Human approval before any send	Campaign studio review step — no auto-send without explicit approval
Negative tag requires written reason	Tag picker enforces note before saving negative tags

Double protection for negative-tagged guests Guests bearing Do Not Contact, Do Not Promote or Unwanted Guest tags are excluded at two independent points: first by the segment filter (they never enter the audience) and second by the consent guardrail layer (they are blocked at send time). A bug in one layer cannot cause an unwanted send.

15 Delivery roadmap

From prototype to production

Phase	Scope	Status
Phase 0 — Prototype	Full HTML/JS prototype: dashboard, guests, segments, campaign studio, deals, automation, consent. Scoring engine and tag manager included.	Done
Phase 1 — API integration	Replace fixture layer with real Chekin guest and reservation APIs. Connect consent service. Deploy behind feature flag.	Next
Phase 2 — Campaign execution	Wire campaign studio to Automation Engine. Real send via email and WhatsApp channels. Analytics tracking.	Planned
Phase 3 — App Sales integration	Live deal inventory from App Sales. Dynamic deal scoring. Revenue attribution.	Planned
Phase 4 — Scaling	Multi-property support. Bulk tagging. Tag import from PMS. Score history and trend.	Future

16 Open questions and decisions

Items requiring alignment before Phase 1

- ✓ **Score persistence** Should the Opportunity Score be computed on every page load (current approach: pure function, no storage) or persisted in a database for historical trending and performance queries?
- ✓ **Tag ownership** Are tags per-property or per-guest-global? A guest staying at two properties might have a VIP tag at one and a Needs Review tag at the other.
- ✓ **Scoring signal weights** The current weights (+25 upcoming, +15 repeat, etc.) are a reasoned first draft. Once the prototype is live with real data, weights should be validated against actual conversion outcomes.
- ✓ **Automation Engine integration** Which triggers does the Automation Engine expose? Are they event-based (check-in completed) or schedule-based (3 days before arrival)? This determines campaign timing fidelity.
- ✓ **Tag audit trail** Should tag assignment history be stored (who added a Whale tag, when, with what note)? Required for compliance if negative tags are added.

17 Appendix: data model and implementation notes

Reference for engineering

Guest profile fields (subset)

GUEST MODEL

```
{
  id:              string,         // guest_001 … guest_029
  fullName:        string,
  email:           string,
  lifecycleStage:  'prospect' | 'guest' | 'in_house' | 'repeat' | 'vip' |
'past',
  totalRevenue:    number,          // EUR
  nextStayAt:      ISO8601 | null,
  lastStayAt:      ISO8601 | null,
  consentSummary:  { email, whatsapp, globalUnsub },
  profileComplete: number,          // 0–100 %
  appSalesPurchases: number,
  tags:            string[],        // booking-derived: 'family', 'couple',
'business'…
  segmentTags:     string[],        // CRM decisions: 'whale', 'vip',
'do_not_contact'…
  segmentTagNotes: Record<string, string>  // note required for negative tags
}
```

Opportunity Score — scoring function signature

SCORING ENGINE

```
computeOpportunityScore(guestId) → {
  score:    number,      // 0–100, capped
  reasons:  string[],   // up to 5 plain-language reasons shown in UI
  tier:     'high' | 'medium' | 'low',
  excluded: boolean     // true if a negative tag forces score to 0
}
```

Segment filter extensions

Filter key	Type	Description
segmentTag	string	Match guests where segmentTags includes this tag slug.
minScore	number	Match guests with opportunityScore ≥ minScore.
excludeNegative	boolean	Exclude guests bearing any negative tag (default: true in campaigns).
search	string	Full-name, email or phone substring match.
lifecycleStage	string	Filter by lifecycle stage.
consent	'marketing' | 'no_marketing'	Filter by email consent state.
upcoming	boolean	Filter guests with a check-in in the next 90 days.